

Abstract

We present Zodiacal, a pure-Rust implementation of blind astrometric plate solving that achieves 98.3% solve rates on simulated narrow-field imagery with a median solve time of 1.42 s. Zodiacal implements the geometric hashing framework of [Lang et al. \(2010\)](#), constructing scale-invariant four-star asterism codes and matching them against prebuilt index files via k -d tree search. Our primary contribution is a systematic analysis of *index sharding strategies*—the partitioning of the quad code space into narrow angular-scale bands—and their effect on solver performance. We show that splitting a single broadband index (10–600″) into 12 narrow bands at a $\sqrt{2}$ scale factor increases the solve rate from 96.0% to 98.3% compared to 4 wider bands at a $3\times$ factor, while reducing the mean verification count by 33%. The improvement arises from smaller k -d trees per band that reduce false-positive code matches, combined with per-index scale filtering that skips irrelevant bands entirely. We characterize the remaining 1.7% failure population as quad-matching timeouts concentrated at moderate Galactic latitudes ($|b| \approx 20\text{--}45^\circ$), where code-space collisions create a false-match flood. The system is fully open-source and designed for integration into automated survey pipelines.

Zodiacal: Accelerating Blind Astrometric Plate Solving Through Multi-Scale Index Sharding

Matthew Goodman
Exclosure
matt@exclosure.io

Keywords: astrometry — methods: data analysis — techniques: image processing — catalogs

1 Introduction

Blind astrometric plate solving—the determination of a celestial image’s pointing, rotation, and plate scale from the star pattern alone—is a foundational capability for astronomical survey pipelines, robotic telescopes, and spacecraft attitude determination. The problem is computationally challenging: given N detected sources in an image and a catalog of $\sim 10^8$ reference stars, one must identify the correct correspondence without prior knowledge of the field’s location on the sky.

The dominant approach, introduced by [Lang et al. \(2010\)](#) and implemented in the widely-used `astrometry.net` system, reformulates the problem as geometric hashing. Four-star asterisms (“quads”) are encoded into a low-dimensional code space that is invariant to rotation, scale, and parity. A k -d tree indexes the code space, enabling sub-linear lookup of candidate matches. Bayesian verification then scores each candidate World Coordinate System (WCS) against the full star field.

While `astrometry.net` has proven highly effective—reporting $> 99.9\%$ solve rates on Sloan Digital Sky Survey (SDSS) fields ([Lang et al., 2010](#))—its implementation is distributed across C, Python, and multiple FITS-based intermediate file formats, making it difficult to embed in modern compiled pipelines. More importantly, the interaction between index construction parameters and solver performance has not been systematically characterized. In particular, the choice of how to *shard* the angular scale space into multiple index files—how many bands, how wide, and at what scale factor—has a first-order effect on solve speed and reliability that is not well-documented in the literature.

In this paper we present Zodiacal, a pure-Rust

reimplementation of the geometric hashing approach that achieves comparable solve rates while enabling systematic experimentation with index construction. Our primary contribution is an empirical analysis of index sharding strategies and their performance implications. We demonstrate that:

1. Narrower angular-scale bands reduce the effective k -d tree size per index, cutting false-positive matches and improving solve speed.
2. Per-index scale filtering—skipping indexes whose band cannot overlap with the inferred backbone angular scale—provides an efficient pruning mechanism that scales with the number of bands.
3. The combination of these effects yields a Pareto improvement: finer sharding increases both solve rate and reduces median solve time for the majority of fields.
4. The residual failure population ($\sim 1.7\%$) is dominated by code-space collisions at moderate Galactic latitudes, a regime where narrower bands provide diminishing returns.

Section 2 describes the solver architecture. Section 3 details the index construction pipeline and sharding strategies. Section 4 presents benchmark results across configurations. Section 5 analyzes the failure population. Section 6 discusses implications and future work. Section 7 extends the scale-only sharding of §3 along a second, spatial axis (the `.zdc1.bundle` format) and characterises region-restricted solving against a depth-9, $G \leq 20$ Gaia bundle. Section 8 describes how proper motion is applied at WCS fit and verification so that out-of-epoch plates (e.g. archival Hubble frames) solve against the same Gaia-epoch bundle.

2 Algorithm

Zodiacal implements the full pipeline described by [Lang et al. \(2010\)](#): source extraction, quad formation, code matching, WCS fitting, and Bayesian verification. We summarize each stage, noting implementation choices that differ from `astrometry.net`.

2.1 Source Extraction

Sources are detected using a DAO/IRAF-style algorithm provided by the `starfield` library ([Goodman, 2025](#)): Gaussian convolution for background estimation, connected-component labeling above a signal-to-noise threshold, and flux-weighted centroiding. The result is an ordered list of (x, y, f) tuples sorted by decreasing flux f .

For pre-extracted source lists (e.g., from upstream pipelines), Zodiacal accepts a JSON interchange format containing pixel positions, fluxes, and optional metadata (known RA/Dec, plate scale hints) that can constrain the search space.

2.2 Quad Codec

A quad consists of four stars (A, B, C, D) . Stars A and B form the *backbone*—the pair with the largest angular separation among the four. The remaining stars C and D are encoded relative to the backbone in a rotation- and scale-invariant coordinate system.

Specifically, let $\mathbf{M}$ be the midpoint of A and B on the unit sphere. All four stars are projected onto the tangent plane at $\mathbf{M}$ via gnomonic projection. A rotation and scaling are applied to normalize the baseline $\overline{AB}$ to unit length along the x -axis, with A at the origin. The positions of C and D in this frame yield a four-dimensional code:

$$\mathbf{c} = (c_x, c_y, d_x, d_y) \in \mathbb{R}^4. \quad (1)$$

Two canonicalization invariants ensure that equivalent geometric configurations produce identical codes regardless of labeling:

1. **Mean- x constraint:** if the mean of the x -coordinates exceeds 0.5, swap $A \leftrightarrow B$ and transform $c_i \leftarrow 1 - c_i$ for all i .
2. **Ordering constraint:** sort C and D by their x -coordinates so that $c_x \leq d_x$.

This codec is identical to the one described in [Lang et al. \(2010\)](#) and implemented in `astrometry.net`'s `quad-utils.c`.

2.3 Two-Phase Quad Generation

Field quads are generated incrementally from the N brightest sources (default $N = 50$), following the two-phase strategy from `astrometry.net`. When star n is introduced:

- **Phase 1:** Star n serves as backbone star B . For each $A < n$, candidates C and D are drawn from $\{0, \dots, n-1\}$ within the bounding circle of the AB pair.
- **Phase 2:** Star n serves as off-diagonal star C . For each backbone pair $A < B < n$, candidate D is drawn from $\{0, \dots, n-1\}$.

This ensures each unique quad is attempted exactly once across all iterations, maximizing search coverage without redundancy. A bounding circle constraint (§2.4) rejects geometrically implausible C and D candidates early.

2.4 The AB-Diameter Inscription Test

The Lang quad invariant requires the two off-diagonal stars C and D to lie inside the *AB-diameter circle*—the circle whose diameter is the segment $\overline{AB}$. In the canonical code-space coordinates of §2.2 (with A at the origin and B at $(1, 0)$), this is the circle of radius $\frac{1}{2}$ centred on $(\frac{1}{2}, \frac{1}{2})$:

$$(x - \frac{1}{2})^2 + (y - \frac{1}{2})^2 \leq \frac{1}{4}, \quad (2)$$

matching `astrometry.net`'s `solver/quad-builder.c::check_inbox`. During quad *construction* the test is evaluated on the sphere before projection, to skip the gnomonic step for candidates that will fail the invariant. Let $\mathbf{M}$ be the midpoint of A and B on the unit sphere and θ_{AB} the angular separation of A and B . The spherical-geometry form of Eq. 2 is

$$\|\mathbf{C} - \mathbf{M}\|^2 \leq 2 \left(1 - \cos \frac{\theta_{AB}}{2} \right), \quad (3)$$

expressed in the squared Euclidean (chord) distance between unit vectors. The right-hand side is the chord² corresponding to angular distance $\theta_{AB}/2$ from $\mathbf{M}$, i.e. the radius of the inscription circle on the sphere.

A subtle error in our original implementation used $2(1 - \cos \theta_{AB})$ for the threshold—the chord² for angular distance θ_{AB} , twice the correct radius—admitting candidates anywhere within distance $|\overline{AB}|$ of $\mathbf{M}$.

rather than $|\overline{AB}|/2$. The encoding-side canonicalisation (mean- x and ordering constraints of §2.2) was unaffected, so codes themselves remained well-formed, but the index admitted quads whose C or D projected outside the validity circle in code space. The bug was caught visually on a Hubble Space Telescope scene render where D sat at 122.7% of the inscription radius despite a correct solve, and fixed by replacing the threshold with Eq. 3 at all three quad-build sites. Bundles rebuilt against the corrected threshold are 20–30% smaller per band, with proportionally fewer spurious code matches per solve; older bundles continue to work unchanged because the verifier already rejects the laxer quads as verification failures.

2.5 Code Matching via k -d Tree

Each prebuilt index file contains a k -d tree over the four-dimensional code space of its constituent quads. Given a field code $\mathbf{c}_f$, the solver performs a range search:

$$\mathcal{M} = \{\mathbf{c}_i \in \mathcal{T} : \|\mathbf{c}_f - \mathbf{c}_i\|^2 < \epsilon\}, \quad (4)$$

where $\mathcal{T}$ is the code tree and $\epsilon = 0.01$ is the default squared Euclidean tolerance. Each match yields a star correspondence $(A_f, B_f, C_f, D_f) \leftrightarrow (A_i, B_i, C_i, D_i)$.

To handle image parity (flipped fields), each field code is tried in both the original and x/y -swapped orientation, doubling the match attempts without doubling the index size.

2.6 WCS Fitting

From each four-star correspondence, a tangent-plane (TAN) WCS is fit via least-squares. The tangent point is the centroid of the matched reference stars’ unit vectors, and the CD matrix is solved from the 2×2 normal equations relating gnomonic-projected sky coordinates to pixel coordinates. Solutions with a singular or near-singular matrix ($|\det| < 10^{-30}$) are rejected. A residual check ensures all four quad stars project to within 10 pixels of their detected positions.

2.7 Bayesian Verification

Following Lang et al. (2010), each candidate WCS is scored using a Bayesian log-odds framework. Reference stars within the projected field of view are identified via a 3D k -d tree range search, then projected to pixel coordinates. A 2D k -d tree over the projected positions enables efficient nearest-neighbor queries.

For each field source s_i , the log-odds contribution is:

$$\Delta_i = \max \left(\log \frac{p_{\text{fg}}(s_i)}{p_{\text{bg}}}, \log \frac{p_{\text{dist}}(s_i)}{p_{\text{bg}}} \right) - \log p_{\text{bg}}, \quad (5)$$

where p_{fg} is a Gaussian likelihood for a genuine match (with $\sigma = r_{\text{match}}/2$, default $r_{\text{match}} = 5$ px), p_{bg} is a uniform background rate, and p_{dist} models a dynamic distractor rate that adapts as matches accumulate.

A solution is accepted when the cumulative log-odds exceeds 20 (odds ratio $\sim e^{20} \approx 5 \times 10^8:1$) with at least 10 matched stars. Early termination at log-odds < -20 avoids wasting time on clearly spurious candidates.

3 Index Construction and Sharding

The quality and structure of prebuilt index files have a dominant effect on solver performance. This section describes the index construction pipeline and introduces the sharding strategies we evaluate.

3.1 HEALPix Uniformization

Raw star catalogs (e.g., Gaia DR3) exhibit orders-of-magnitude variation in stellar density between the Galactic plane and poles. Naive quad generation from such catalogs concentrates quads in dense regions, leaving sparse regions underrepresented.

We uniformize the catalog using HEALPix (Górski et al., 2005) pixelization in the nested indexing scheme. The sky is divided into $12 \times N_{\text{side}}^2$ equal-area cells, where $N_{\text{side}} = 2^d$ at depth d . The depth is automatically computed from the index’s upper scale bound:

$$d = \left\lceil \log_2 \frac{\sqrt{\pi/3}}{\theta_{\text{max}}} \right\rceil, \quad (6)$$

where θ_{max} is twice the upper angular scale (in radians), capped at $d \leq 8$ (786,432 cells) for memory efficiency.

For each cell, a max-heap retains only the k brightest stars (default $k = 30$), producing a spatially uniform subset of the input catalog. This is a streaming single-pass algorithm over the catalog.

3.2 Per-Cell Quad Building

Quads are generated per HEALPix cell (at a potentially different depth than the uniformization depth). For each cell, we collect stars from the cell and its 8

neighbors (to avoid boundary effects), build a 3D k -d tree, and generate quads whose backbone angular distance falls within $[\theta_{\min}, \theta_{\max}]$.

Key parameters controlling quad density:

- **Passes** (default 16): number of quad-building sweeps per cell.
- **Max reuse** (default 8): maximum number of quads in which any single star may appear, enforced via atomic counters.
- **Max quads**: global cap, distributed uniformly across cells.

Quad generation is parallelized across cells using `rayon` with an atomic done-flag for early termination when the global quota is met.

3.3 Sky-Uniform Coverage at Depth

Per-cell quad building (§3.2) achieves spatial uniformity *by design*, but only if the iteration order of the construction loop respects the cell partition. Our initial production builder (`src/index/builder.rs`) sorted the entire post-uniformization catalog by brightness and emitted quads anchor-by-anchor under a single global `max_quads` cap. In dense regions a single bright anchor can emit $\mathcal{O}(K^2)$ quads with $K \sim 10^2$ neighbors, so the global cap saturates within the first few thousand anchors and faint anchors in sparse cells never get visited. The empirical cost is severe: on our `gaia_g14.zdc1` production index (1M quads, 16.8M stars at $G \leq 14$), only **133 of 12,288** HEALPix-5 cells (1.1%) hosted any quads, and a single $G = 7.6$ star participated in $\sim 2.6 \times 10^5$ quads—roughly a quarter of the entire budget. A 0.5° FOV at random sky pointings drew zero quads in 99.8% of trials.

Two coupled changes restore sky uniformity. First, the cell-driven builder (`src/index/cell_builder.rs`) walks $0..N_{\text{cells}}$, pulls per-cell stars from a level-5-sharded input source, brightness-truncates to `max_stars_per_cell`, and emits up to `quads_per_cell` quads using only that cell’s stars. The per-cell budget makes uniformity a *structural* property: every populated cell receives an independent quota.

Second, the magnitude limit must be deep enough that the per-cell budget—not catalog depth—is the binding constraint. At Gaia $G \leq 14$ the median high-galactic-latitude cell holds $\sim 1,400$ stars, which can geometrically support ~ 100 valid quads only marginally. At $G \leq 20$ every populated cell holds $\sim 8.5 \times 10^4$ stars and the per-cell brightness truncation kicks in well before the cell’s star geometry runs

out. Going deeper is now free of the previous “deeper $\Rightarrow$ even more bright-region quads” pathology because the cell loop bounds per-cell quad emission directly.

The deep build pipeline rests on three pieces of infrastructure not needed by the in-RAM builder:

- **Streaming sidecar.** The Gaia refinement sidecar (§2.6) is the per-quad memory hog at ~ 88 B/record. The streaming writer (`src/refinement/sidecar.rs::SidecarStreamWriter`) accepts per-cell chunks, sorts each in RAM to a numbered temp file, and runs an external k -way merge at finalize time. Peak RAM is bounded by the largest per-cell chunk ($\sim$ tens of MB at level 5). Output bytes are identical to the in-RAM writer for the same input set.
- **Resumable manifest.** A JSON manifest (`src/index/build_manifest.rs`) tracks completed cells and is atomically rewritten after each per-cell commit; per-cell artifacts live under a work directory keyed by `cell_id`. A re-run with the same `--work-dir` skips already-committed cells.
- **Order-stable parallelism.** The cell loop runs under `rayon::par_iter`; `finalize` iterates the manifest’s `BTreeSet<u32>` of completed cells (sorted by `cell_id` regardless of completion order) and the sidecar merge sorts chunks by path before merging. A dedicated test (`parallel_build_matches_sequential`) asserts that 1-thread and 4-thread builds produce byte-identical `.zdc1` and `.zdc1.gaia` files.

The remaining failure mode is geometric: a high-latitude void with too few faint stars cannot fill its `quads_per_cell` budget no matter how the loop is structured. The build summary surfaces per-cell counts so this is observable. Quads whose backbones straddle a HEALPix-5 boundary at the source’s depth are also dropped under the strict per-cell convention; at the typical quad scales of $30''$ to $30'$ this loss is small relative to the cell width ($\sim 1.8^\circ$) but is nonzero. Adding neighbor context to the per-cell builder is a follow-up.

3.4 Sharding Strategies

A *sharding strategy* partitions the full angular scale range $[\theta_{\min}, \theta_{\max}]$ into B bands, each producing an independent index file. Each band b covers the range $[\theta_b^{\text{lo}}, \theta_b^{\text{hi}}]$ where:

$$\theta_b^{\text{hi}} = \min(\theta_b^{\text{lo}} \cdot \alpha, \theta_{\max}), \quad (7)$$

and α is the *scale factor* controlling band width. The number of bands is:

$$B = \left\lceil \frac{\log(\theta_{\max}/\theta_{\min})}{\log \alpha} \right\rceil. \quad (8)$$

We evaluate three strategies spanning a range of granularity:

1. **Monolithic** ($B = 1$): a single index covering 10–600". This serves as the baseline.
2. **Coarse sharding** ($\alpha = 3$, $B = 4$): bands at 10–30, 30–90, 90–270, 270–600".
3. **Fine sharding** ($\alpha = \sqrt{2}$, $B = 12$): bands at approximately octave intervals (10–14, 14–20, ..., 425–600").

3.4.1 Per-Index Scale Filtering

When multiple narrow-band indexes are loaded, the solver exploits the known scale bounds to skip irrelevant indexes for each field quad. Given a backbone pixel distance d_{AB} and the pixel scale range $[s_{\min}, s_{\max}]$ (in rad/px), the inferred angular scale range of the backbone is:

$$[\theta^{\text{lo}}, \theta^{\text{hi}}] = [d_{AB} \cdot s_{\min}, d_{AB} \cdot s_{\max}]. \quad (9)$$

An index with band $[\theta_b^{\text{lo}}, \theta_b^{\text{hi}}]$ is skipped if these intervals do not overlap:

$$\theta^{\text{lo}} > \theta_b^{\text{hi}} \quad \text{or} \quad \theta^{\text{hi}} < \theta_b^{\text{lo}}. \quad (10)$$

This is evaluated in $O(B)$ per quad attempt and has negligible overhead. With fine sharding, a typical field quad overlaps at most 2–3 of the 12 bands, eliminating ~75–85% of code tree searches.

3.4.2 Why Narrower Bands Improve Performance

The performance benefit of narrow bands arises from two complementary mechanisms:

Smaller code trees. A monolithic index spanning 10–600" contains all quads across the full scale range. A narrow band (e.g., 90–127") contains only quads whose backbone falls within that octave. Assuming roughly uniform quad density across log-scale, each band's k -d tree is $\sim B$ times smaller, reducing the expected number of false-positive code matches per range search by the same factor.

Concretely, in our Gaia magnitude-19 experiments, the monolithic index contained ~58M quads, while the 4-band configuration distributed these as 15K, 1.2M, 12.4M, and 12.5M quads per band.

Scale filtering. The per-index filtering described in §3.4.1 avoids searching the code tree altogether for out-of-band indexes. This amortizes the cost of loading multiple indexes: although B indexes are loaded, only ~2–3 are searched per quad, so the aggregate search cost is lower than searching a single monolithic tree.

The combination yields a Pareto improvement: fine sharding increases the solve rate *and* reduces median solve time for the typical case.

3.5 The k -d Forest: Multi-Index Search

When B narrow-band indexes are loaded simultaneously, the solver searches a k -d *forest*—a collection of B independent four-dimensional k -d trees, each covering a different angular scale band. We derive the analytical scaling of this architecture and its interaction with plate scale knowledge.

3.5.1 False-Positive Rate in the Code Space

The quad code $\mathbf{c} \in \mathbb{R}^4$ (Eq. 1) is scale- and rotation-invariant, so codes from all angular scales populate the same four-dimensional space. A range search with squared L^2 tolerance ϵ (default 0.01) queries a 4-ball of radius $r = \sqrt{\epsilon}$, whose volume is

$$V_\epsilon = \frac{\pi^2}{2} r^4 = \frac{\pi^2}{2} \epsilon^2. \quad (11)$$

For $\epsilon = 0.01$, this gives $V_\epsilon \approx 4.93 \times 10^{-4}$.

Let $\rho(\mathbf{c})$ be the local density of codes in the space. The expected number of matches from a tree containing Q codes is

$$\langle M \rangle = \int_{\|\mathbf{c} - \mathbf{c}_f\|^2 < \epsilon} \rho(\mathbf{c}) d\mathbf{c} \approx \rho(\mathbf{c}_f) \cdot V_\epsilon, \quad (12)$$

where $\mathbf{c}_f$ is the field code and the approximation holds when ρ is locally smooth on the scale of r . For a uniform distribution over unit volume, $\rho = Q$ and $\langle M \rangle \approx Q \cdot V_\epsilon$.

The key observation is that the codes are scale-invariant, but different angular scale bands sample *different stellar neighborhoods* on the sky, producing codes with *independent* false-positive populations. A code that happens to lie near $\mathbf{c}_f$ in band b is a coincidental geometric similarity between unrelated star patterns at scale θ_b —it has no correlation with near-matches at scale $\theta_{b'}$.

3.5.2 Partitioning the False-Positive Budget

For a monolithic index of $Q = \sum_{b=1}^B Q_b$ total quads with density $\rho_{\text{mono}} = \sum_b \rho_b$, the expected false

matches per query are

$$\langle M_{\text{mono}} \rangle \approx \rho_{\text{mono}}(\mathbf{c}_f) \cdot V_\epsilon = \sum_{b=1}^B \rho_b(\mathbf{c}_f) \cdot V_\epsilon. \quad (13)$$

When searching the forest with scale filtering, only B_{eff} bands overlap the backbone’s angular scale interval. The false matches from the searched bands are

$$\langle M_{\text{forest}} \rangle = \sum_{b \in \mathcal{S}} \rho_b(\mathbf{c}_f) \cdot V_\epsilon, \quad (14)$$

where $\mathcal{S}$ is the set of overlapping bands with $|\mathcal{S}| = B_{\text{eff}}$.

If the quads are distributed approximately uniformly across bands ($Q_b \approx Q/B$, hence $\rho_b \approx \rho_{\text{mono}}/B$), the false-positive reduction factor is

$$\boxed{\frac{\langle M_{\text{forest}} \rangle}{\langle M_{\text{mono}} \rangle} = \frac{B_{\text{eff}}}{B}} \quad (15)$$

This is the central result: *the k -d forest reduces false positives by the ratio of searched bands to total bands.* For $B = 12$ and $B_{\text{eff}} = 2$, this is a $6\times$ reduction.

3.5.3 Effective Band Overlap as a Function of Scale Bracket

The number of overlapping bands B_{eff} depends on the relationship between the plate scale bracket and the band width. Given a backbone pixel distance d_{AB} and plate scale bracket $[s(1 - \delta), s(1 + \delta)]$ (where $\delta = \Delta s/2s$ is the fractional half-width), the backbone angular scale spans

$$\theta \in [d_{AB} \cdot s(1 - \delta), d_{AB} \cdot s(1 + \delta)]. \quad (16)$$

For geometric bands with scale factor α , each band spans a factor α in angular scale. The angular bracket spans a factor $(1 + \delta)/(1 - \delta) \approx 1 + 2\delta$ for small δ . The number of bands overlapped by this interval is

$$B_{\text{eff}} = 1 + \left\lceil \frac{\log\left(\frac{1+\delta}{1-\delta}\right)}{\log \alpha} \right\rceil. \quad (17)$$

Table 1 evaluates this for representative configurations.

With coarse $3\times$ bands, even moderate brackets fit within a single band, so $B_{\text{eff}} = 1$ and the reduction is exactly B . With fine $\sqrt{2}$ bands, the bracket may span 2-3 bands, but there are $3\times$ more bands total, so the net reduction is comparable or better.

Table 1: Effective bands searched B_{eff} and false-positive reduction factor B/B_{eff} for different scale brackets and band widths.

Scenario	δ	$\alpha = 3$	$\alpha = \sqrt{2}$
		B_{eff} (red.)	B_{eff} (red.)
Known ($\pm 5\%$)	0.05	1 ($4\times$)	1-2 ($6\text{--}12\times$)
Moderate ($\pm 20\%$)	0.20	1 ($4\times$)	2-3 ($4\text{--}6\times$)
Coarse ($\pm 50\%$)	0.50	1-2 ($2\text{--}4\times$)	3-4 ($3\text{--}4\times$)
Unknown ($\delta \rightarrow \infty$)	—	4 ($1\times$)	12 ($1\times$)

3.5.4 Impact on Solve Time

The total solver time is dominated by verification cost (WCS fitting plus Bayesian scoring) applied to each code match. Let C_v denote the per-verification cost and F the number of field quads attempted. Then

$$T_{\text{solve}} \approx F \cdot \langle M \rangle \cdot C_v + T_{\text{overhead}}, \quad (18)$$

where T_{overhead} includes quad generation, code computation, and k -d tree traversal (all subdominant).

Substituting Eq. 15:

$$\frac{T_{\text{mono}}}{T_{\text{forest}}} \approx \frac{\langle M_{\text{mono}} \rangle}{\langle M_{\text{forest}} \rangle} = \frac{B}{B_{\text{eff}}}. \quad (19)$$

The expected speedup from the k -d forest is B/B_{eff} . For our configuration ($B = 12$, $B_{\text{eff}} \approx 2$), this predicts a $\sim 6\times$ reduction in the number of false-positive verifications.

This prediction is consistent with the observed data: failed cases under the 4-band configuration verify a median of 259,711 candidates vs. a median of 3,628 for solved cases. The ratio of $\sim 70\times$ between failed and solved cases far exceeds the B/B_{eff} speedup, confirming that the failures represent a qualitatively different regime (code-space degeneracy) rather than merely a quantitative slowdown.

3.5.5 k -d Tree Traversal Scaling

Beyond reducing false positives, the forest also reduces k -d tree traversal cost. For a balanced k -d tree in d dimensions, the worst-case range search cost is $O(Q^{1-1/d} + M)$ (Lee and Wong, 1977). In $d = 4$ dimensions:

$$C_{\text{tree}}(Q) = O(Q^{3/4} + M). \quad (20)$$

For the forest, the aggregate traversal cost across B_{eff} bands is

$$C_{\text{forest}} = \sum_{b \in \mathcal{S}} O(Q_b^{3/4} + M_b) = B_{\text{eff}} \cdot O\left(\left(\frac{Q}{B}\right)^{3/4} + \frac{M_{\text{mono}}}{B}\right). \quad (21)$$

The traversal speedup is

$$\frac{C_{\text{mono}}}{C_{\text{forest}}} = \frac{Q^{3/4} + M_{\text{mono}}}{B_{\text{eff}} ((Q/B)^{3/4} + M_{\text{mono}}/B)} = \frac{Q^{3/4}}{B_{\text{eff}} Q^{3/4}/B^{3/4} + B_{\text{eff}} M_{\text{mono}}/B} \quad (22)$$

In the match-dominated regime ($M_{\text{mono}} \gg Q^{3/4}$), this reduces to B/B_{eff} as before. In the traversal-dominated regime ($M \ll Q^{3/4}$), the speedup is $B^{3/4}/B_{\text{eff}}$, which for $B = 12$, $B_{\text{eff}} = 2$ gives $12^{3/4}/2 \approx 3.2\times$.

In practice, range searches with $\epsilon = 0.01$ return enough matches that the match-dominated regime applies for large indexes ($Q > 10^6$), making the B/B_{eff} approximation accurate.

3.5.6 Scaling with FoV Bracket Width

The false-positive reduction (Eq. 15) depends on B_{eff} , which in turn depends on how tightly the plate scale is known. We summarize three operational regimes:

Known plate scale ($\delta < (\alpha - 1)/2$). The bracket fits within one band width. $B_{\text{eff}} = 1-2$, and the forest provides the full $B/B_{\text{eff}} \approx B$ reduction. This is the common case for survey pipelines with calibrated optics.

Unknown plate scale ($\delta \rightarrow \infty$). All bands are searched ($B_{\text{eff}} = B$), and the forest provides no false-positive reduction over monolithic. However, it imposes no *penalty* either: the total work is $\sum_b (\log Q_b + M_b) \approx \log Q + M_{\text{mono}}$, since the matches partition across bands with negligible root-traversal overhead.

Intermediate bracket ($\delta \sim 1$). The bracket spans $\sim B/3$ bands, giving a modest $\sim 3\times$ reduction. This is the regime where the choice of α matters: finer bands ($\alpha = \sqrt{2}$) yield more bands and hence a larger B/B_{eff} ratio, despite each bracket spanning more bands.

3.5.7 Memory and Loading Trade-offs

The k -d forest partitions the same total quad population across B trees, so aggregate memory is approximately constant:

$$\mathcal{M}_{\text{forest}} \approx \mathcal{M}_{\text{mono}} + B \cdot \mathcal{O}_{\text{root}}, \quad (23)$$

where $\mathcal{O}_{\text{root}}$ is the per-tree overhead (root node, metadata, allocation headers), typically < 1 KB. For our 12-band configuration, total load time is ~ 15 s and peak memory is ~ 5 GB—comparable to a monolithic index of the same quad count.

4 Results

4.1 Test Configuration

We evaluate on 1,000 simulated stellar fields generated by `meter-sim`:

- Image dimensions: 9568×6380 px
- Plate scale: $0.1296''/\text{px}$
- Field of view: $\sim 20' \times 13'$
- Timeout: 60 s per field

Source lists were pre-extracted and provided as JSON, bypassing the extraction step to isolate solver performance. All runs used a single machine with rayon parallelism disabled for timing consistency.

Two Gaia DR3 catalogs were used:

- Magnitude ≤ 16 : ~ 83 M stars (2.5 GB)
- Magnitude ≤ 19 : ~ 565 M stars (17 GB)

4.2 Sharding Comparison

Table 2 summarizes the solver performance across index configurations.

The key findings are:

Finer sharding improves solve rate. Moving from 4 bands ($\alpha = 3$) to 12 bands ($\alpha = \sqrt{2}$) increases the solve rate from 96.0% to 98.3%, recovering 23 additional fields. These fields had been lost to false-positive floods in the wider bands.

Finer sharding reduces verification load. The mean number of WCS verifications per image drops from 21,301 under 4-band sharding to 14,139 under 12-band sharding—a 33% reduction consistent with the analytical prediction of B_{eff}/B fewer false positives per searched band (§3.5). The forest searches fewer, smaller trees, reducing both the number of false-positive code matches and the associated verification cost.

4.3 Solve Time Distribution

For the best configuration (magnitude 19, 12 bands), the cumulative solve time distribution is:

Figures 1 and 2 show the solve time distributions for the 4-band and 12-band configurations respectively, and Figure 3 shows the corresponding cumulative distribution functions.

The median solve time of 1.42 s under 12-band sharding is comparable to the 1.05 s achieved by

Table 2: Solver performance across index configurations on 1,000 simulated fields. All configurations use the same underlying catalog and solver parameters; only the index sharding strategy differs.

Configuration	B	α	Solved	Rate	Median t	Mean verif.
Mag 16, 4 bands	4	3.0	960	96.0%	1.05 s	21,301
Mag 19, 12 bands	12	$\sqrt{2}$	983	98.3%	1.42 s	14,139

Table 3: Cumulative solve time distribution for the 983 solved fields under the best configuration (Mag 19, 12 bands, $\alpha = \sqrt{2}$).

Time threshold	Fraction solved
< 1 s	44%
< 5 s	78%
< 10 s	89%
< 30 s	98%
< 60 s	100%




Figure 1: Solve time histogram for the 4-band configuration ($\alpha = 3$). Timeouts at 60 s are shown in red. The dashed line marks the median solve time.

4-band sharding, despite searching a larger aggregate index. The 12-band configuration performs 33% fewer verifications per image (14,139 vs. 21,301), confirming the k -d forest’s false-positive reduction (§3.5). The slight increase in median time reflects a handful of harder fields that benefit from the finer bands but take longer to converge. This trade-off is offset by a 2.3 percentage-point improvement in solve rate.

4.4 Verified Candidate Count

The number of WCS candidates verified before finding the correct solution is a direct measure of the false-positive rate in code matching. Table 4 contrasts this metric between solved and failed cases.

Under both configurations, failed cases evaluate




Figure 2: Solve time histogram for the 12-band configuration ($\alpha = \sqrt{2}$). The timeout population shrinks from 40 to 17 fields.

Table 4: Number of WCS candidates verified during a solve attempt.

Config	Outcome	Min	Median	Max
4-band	Solved	13	3,596	471,309
4-band	Failed	133,533	226,861	325,582
12-band	Solved	50	4,440	268,884
12-band	Failed	17,012	159,399	210,378

40–60× more candidates than the median solved case. Notably, the 12-band forest reduces the median failure verification count from 226,861 to 159,399—a 30% reduction—consistent with the B/B_{eff} scaling derived in §3.5. This confirms that failures are driven by false-positive code matches flooding the solver’s search budget, and that narrower bands partially mitigate this even in the degenerate regime.

5 Failure Analysis

We investigated the 17 remaining failures (1.7%) using the Gaia DR3 catalog (82.7M stars to magnitude 16) to characterize the stellar environment at each pointing.




Figure 3: Cumulative distribution of solve times. The 12-band configuration achieves a higher asymptotic solve rate while maintaining comparable speeds for the majority of fields.

Table 5: Properties of solved vs. failed fields under the 12-band configuration. Stellar density is the Gaia mag-16 star count within a $5^\circ \times 5^\circ$ box centered on the field. Source count is the number of detected sources in the field.

Property	Statistic	Solved	Failed
Source count	Median	200	200
	Mean	164	166
Stellar density (5° box)	Median	16,357	21,641
	Mean	33,605	40,602
$ b $ (deg)	Median	30.0	34.9
	Mean	33.2	30.7
Min NN dist (px)	Median	84.9	111.4
Backbone dist (px)	Median	3,938	5,346

5.1 Failure Population Properties

Table 5 compares the properties of solved and failed fields across several dimensions.

Neither source count, Galactic latitude, nor local stellar density are strong predictors of failure. The failure rate is 2.0% among dense fields (200 sources, the extraction cap) and 1.2% among sparser fields—a marginal difference. Similarly, failure rates are comparable across all stellar density quartiles (0.8–3.2%).

5.2 Distinguishing Characteristics

Two subtle geometric properties distinguish the failure population:

Large backbone distance. Failed fields have a median backbone distance (separation between the two brightest sources) of 5,346 px, 36% larger than the solved median of 3,938 px. A large backbone produces a quad with a correspondingly large angular scale, which—at the test plate scale of $0.1296''/\text{px}$ —maps to $\sim 11.5'$, near the upper end of the index’s scale range. Long backbones sample a larger sky area, increasing the probability of code-space collisions with unrelated star patterns at similar angular scale.

Well-separated sources. The median minimum nearest-neighbor distance among the top 50 sources is 111 px for failed fields vs. 85 px for solved fields. This indicates a more uniform spatial distribution of bright sources, which reduces the diversity of quad geometries available to the solver. Clustered sources produce a richer variety of quad codes, increasing the likelihood that at least one matches uniquely.

5.3 Failure Mechanism

The 17 failures share a common mechanism: the solver evaluates 121,000–210,000 candidate WCS solutions without finding the correct quad match, exhausting the 60 s time budget. The verification count for failed fields is $36\text{--}48\times$ the solved median of 4,440, indicating a false-match flood in the code space.

This is fundamentally a *search budget* problem rather than a verification sensitivity problem. The correct quad may exist in the index but is buried among too many geometrically similar candidates. Longer timeouts could recover some cases, but the flat log-odds plateau at ~ 50 (with only 4 matches) suggests the solver is not converging—it repeatedly finds wrong quads with coincidentally similar codes.

6 Discussion

6.1 Sharding as a First-Order Design Decision

Our results demonstrate that index sharding is not merely an engineering convenience for managing file sizes—it is a first-order algorithmic parameter that directly affects solver accuracy. The 2.3 percentage-point improvement from 4 to 12 bands exceeds the 1.0 percentage-point gain from tripling the catalog depth (83M to 565M stars).

This finding has practical implications: operators can improve solve performance by rebuilding indexes with finer sharding at *zero* cost in catalog data, hardware, or solver code. The only trade-off is increased

disk usage and index load time, both of which are modest (12 files instead of 4, loaded once at startup).

6.2 Diminishing Returns of Band Splitting

The progression from 1 to 4 to 12 bands shows diminishing returns: the largest gains come from the initial split from monolithic to multi-band. Beyond ~ 12 bands, the marginal benefit of splitting further is limited by two factors:

1. At the finest scales (10–14"), HEALPix cells at depth 8 are much larger than the quad angular scale, producing sparse quad populations per cell.
2. The remaining failures are code-space degeneracies that narrower bands cannot resolve—the wrong quads have genuinely similar codes to the correct one.

6.3 Memory and Scale Trade-offs

Fine-scale bands at high HEALPix depth present a memory challenge. Depth 8 corresponds to 786,432 cells, and the k -d tree over a 12M-quad index requires ~ 15 GB of working memory. Depth 9 (3.1M cells) or depth 10 (12.6M cells) would improve quad density for fine bands but requires either streaming construction or significantly more memory. We cap the depth at 8 as a practical compromise.

6.4 Comparison with `astrometry.net`

Direct comparison with `astrometry.net` is complicated by differences in test data (our simulated narrow-field imagery vs. SDSS’s wider-field survey data), index construction parameters, and runtime environments. The reference system reports $>99.9\%$ on SDSS fields (2048×2048 px, $0.396''/\text{px}$, $\sim 13.5'$ FOV), compared to our 98.3% on narrower-field simulations ($\sim 20' \times 13'$ FOV, $0.1296''/\text{px}$).

The narrower field and finer plate scale in our test set present a harder matching problem: fewer reference stars in the field of view, smaller angular scale quads, and tighter tolerances for positional matching. The algorithmic approach is identical; differences in performance are primarily attributable to the test data characteristics.

6.5 Pure-Rust Implementation

Zodiacal is implemented in $\sim 4,000$ lines of Rust with no unsafe code, no C bindings, and no runtime dependencies beyond the standard library and pure-Rust

crates. The const-generic k -d tree (`KdTree<const DIM: usize>`) specializes at compile time for 3D (star positions), 4D (quad codes), and 2D (pixel-space verification), avoiding runtime dispatch overhead.

The monolithic index format (`.zdc1`) bundles metadata, stars, quads, and codes into a single binary file with a JSON metadata header, simplifying distribution compared to `astrometry.net`’s multi-file FITS-based format.

6.6 Future Directions

Several avenues could address the remaining 1.7% failure rate:

- **Finer band splitting in the 30–90" range**, where the bulk of mid-latitude failures concentrate.
- **Higher HEALPix depth** for fine-scale bands, requiring streaming index construction to manage memory. §7 addresses both items above by lifting the single-index design into a per-cell, multi-band bundle that builds at $d_b=9$ on a 64 GB host and is queryable region-by-region; the resulting solve rate against the regenerated test corpus is 100%.
- **Adaptive backbone selection**: instead of always using the brightest stars, choose backbone pairs that maximize geometric diversity.
- **SIP distortion models**: extending beyond TAN to handle optical distortions in wide-field imagers.
- **Proper-motion-aware verification**, addressed in §8: the verifier now propagates each catalog star to the observation epoch when one is supplied, recovering high-PM detections on archival plates whose observation time is far from the Gaia DR3 reference epoch.

7 Spatial Sharding and Region-Restricted Solving

The sharding analysis in §3.4 treated the index as a single global k -d forest, partitioned only along the scale axis. As catalog depth grows toward Gaia DR3’s $G \leq 20$ limit (roughly 9×10^8 stars), this single-file design exceeds working-set sizes that fit comfortably in RAM on conventional hardware: a depth-7 broadband index at $G \leq 20$ alone occupies tens of GB. We therefore introduce a second axis of sharding—spatial—and the supporting on-disk format, build

pipeline, and query path. The combination of scale and spatial sharding lets the solver load only the few thousand cells whose HEALPix footprint overlaps the field of view, rather than mmaping the entire catalog.

7.1 Two-Axis Sharding: The .zdc1.bundle Format

A *bundle* is a directory tree with two parallel children:

- `quads/cell_NNNNN.zqd` — one file per HEALPix cell at a configurable *bundle depth* d_b (we use $d_b = 9$, giving $12 \cdot 4^9 = 3,145,728$ cells of side $\sqrt{4\pi/N} \approx 6.87'$). Each file packs *every* scale band’s quads and codes for that cell into a single 32-byte header plus band-table-indexed payload.
- `gaia/cell_NNNNN.zga` — the matching per-cell Gaia record block (104 bytes per star: `source_id`, ICRS position, full astrometric solution, photometry, and quality flags), enabling the solver’s verification stage to look up catalog stars by index without a second file open.

A top-level `manifest.json` carries band scales, build metadata, and per-band populated-cell counts. The format is documented at the byte level in `docs/bundle-format.md`.

The two-axis layout shifts sharding from one decision per scale band to one decision per (scale band, sky region) pair. Querying then reduces to: enumerate cells overlapping the user-supplied `SkyRegion`, mmap their `.zqd/.zga` entries through a bounded LRU cache, and concatenate the per-band fragments into the `Index` structures the solver consumes.

7.2 Cell-Driven Build Pipeline

Building a depth-9, 7-band, $G \leq 20$ bundle from the Gaia catalog exercises every memory-bound corner of the pipeline. Three structural fixes were necessary to complete on a 64-GB host:

1. **Locality-grouped iteration.** The orchestrator sorts cells by their parent source-cell at the catalog’s native HEALPix depth and groups consecutive same-parent cells into a single rayon work unit. Random per-cell access from 80 worker threads otherwise puts every worker on a different multi-GB parent partition simultaneously and exceeds RAM.
2. **Pre-truncated cached partitions.** The *per bundle-cell* brightness cap N_{stars} is applied during partition build, not after read. For

galactic-plane source cells a single bundle-cell bucket can hold 10^6 stars before the orchestrator’s `truncate(N_stars)` discards 99%; the pre-truncate keeps cached partitions bounded at $\sim 64 \cdot N_{\text{stars}}$ (record size).

3. **Bounded concurrent CSV parses.** A condition-variable permit pool caps simultaneous in-flight catalog parses to a small constant. Each parse transiently holds a multi-GB Gaia row buffer plus an un-truncated bundle-bucket; without the cap, a burst of cache misses across the worker pool stacks transient memory unboundedly.

The build operates in two phases: a parallel *cell phase* that emits the work-directory `.zqd/.zga` shards plus a crash-safe build manifest after each cell, and a *tidy phase* (§7.3) that promotes the populated work directory into a finished bundle. A dedicated *manifest actor* thread owns the `BuildManifest` exclusively and serialises it to disk on a wall-clock cadence, decoupling worker throughput from the manifest fsync rate. With these in place a 7-band, $G \leq 20$, $d_b=9$ bundle ($\sim 9 \times 10^8$ stars, 1.6×10^9 quads) builds end-to-end in roughly 14 hours on an 80-core x86 host; peak resident memory stays under 8 GB across both phases.

7.3 Commit-Edge Atomicity: Eliminating the Tidy Copy

The work directory produced by the cell phase already has the on-disk layout the finished bundle exposes: per-cell shards at `quads/cell_NNNNN.zqd` and `gaia/cell_NNNNN.zga`, named by their HEALPix index. The original tidy phase nevertheless treated the work directory as a staging area: it copied every shard pair into a `<output>.partial/` sibling, `fsync`’d each file, then atomically renamed the sibling onto `<output>`. On a depth-9 $G \leq 20$ bundle that meant copying $\sim 6 \times 10^6$ files and issuing $\sim 6 \times 10^6$ `fsync(2)` calls — a serialised, journal-bound loop that ran about ten hours on the 80-core test bench, comparable in wall-clock to the cell phase itself.

Three progressive simplifications collapse that phase by roughly six orders of magnitude:

1. **Parallelise the copy.** Cell shard pairs land in disjoint paths and never share mutable state, so the loop converts trivially to `populated.par_iter().try_for_each(...)`. Many concurrent `fsync`’s let the kernel coalesce journal commits; throughput rises from one

fsync at a time to the device-sustained sync rate, about a $30\times$ wall-clock reduction in practice.

2. **Drop the per-file fsync.** A bundle is “real” only once its `manifest.json` is present; everything else is cell-shard data that the manifest enumerates. The manifest is written last and itself uses an atomic `.partial.json+rename` pattern, so a crash mid-tidy leaves either no manifest (the bundle does not exist) or a complete one (the bundle does). Per-file `fsync` adds no atomicity guarantee beyond what the manifest already provides; the Linux page cache flushes the shards within seconds either way. Dropping the `fsync` calls yields roughly another $30\times$ reduction.

3. **Rename in place rather than copy.** Once the copy is no longer load-bearing, it has no purpose at all: the work directory’s layout *is* the bundle layout. Tidy reduces to writing `manifest.json` into `work_dir/` and then a single `rename(work_dir, output)`. POSIX `rename(2)` on the same filesystem is atomic against namespace observers: a concurrent reader sees either the pre-tidy or post-tidy state, never a half-merged one. Wall-clock cost: microseconds.

The resulting commit edge is essentially free: bundle finalisation costs the time to serialise the manifest ($\sim$ a few hundred milliseconds at depth 9) plus one `rename(2)` syscall. The end-to-end build of the depth-9 bundle drops from ~ 24 h on the original copy-and-fsync pipeline to the ~ 14 h figure cited above; what remains is the catalog-bound cell phase, which is the lower bound any honest implementation has to pay.

The `tidy_to_zip` entry point still streams from the work directory without consuming it, because a single `ZipWriter` is the natural commit cursor and parallelising it would require a different archive format. The two tidy forms are independent: to emit both representations from one build, run `zip` first (it preserves the work directory) and then `folder` (it consumes it).

7.4 FOV-Aware Choice of Bundle Depth

The bundle depth d_b should be chosen so that a single HEALPix cell is comfortably smaller than the target field of view. At our test camera’s FOV ($20.65' \times 13.77'$), $d_b = 8$ ($13.7'$ cells) is right at the FOV height: cells almost never fit fully inside the




Figure 4: Mean quads per HEALPix cell across the sky for the depth-9, $G \leq 20$ bundle ($N = 3,145,726$ populated cells, total 1.60×10^9 quads). The cap of 560 quads/cell (80×7 bands) is saturated across the entire plane and high-latitude sky; the Galactic poles (NGP and SGP marked) are the only regions where the per-cell quad-builder runs out of candidate stars before reaching the cap. Cell area is computed analytically from `.zqd` file sizes via the relation $n_q = (s - 200)/48$, requiring no catalog access.

image footprint, and the bundle’s `quads_per_cell` selection rarely contains a 4-tuple whose all four stars project onto the visible portion. We saw exactly that failure mode at $d_b = 8$: a triage of the failing cases showed catalog and detected sources in perfect agreement (median residual 1.5 px after WCS projection through truth) but `n_quads_on_image` = 0 for every failure—no findable quad. Going to $d_b = 9$ ($6.87'$) leaves ~ 6 cells overlapping any given FOV, with 2–3 fully inside, and closes that gap entirely.

The corresponding rule of thumb: pick d_b such that $\sqrt{4\pi/(12 \cdot 4^{d_b})} \lesssim \frac{1}{2} \cdot \min(\text{FOV})$. Once d_b is right, lower `max_stars_per_cell` proportionally to the area shrinkage to keep per-cell density invariant: a $4\times$ smaller cell collects $4\times$ fewer brightness-prioritised stars under the same target density. The depth-9 bundle in §7.6 uses `max_stars_per_cell` = 2500, four times smaller than the depth-8 cap.

Figure 4 illustrates the result. Over the sphere, the median quads-per-cell saturates at the configured cap ($80 \times 7 = 560$) everywhere except a small pair of dips at the Galactic poles, where the catalog has too few bright stars to fill the per-band quad budget after `max_reuse` caps. No band-by-band post-processing is required to obtain this uniformity: it falls out of the per-cell quad-builder’s brightness-priority and `max_reuse` bounds.

7.5 Region-Restricted Solving

The query API takes a `SkyRegion(centre, radius)`. `Zdc1Bundle::load_region` enumerates the cells whose centres fall within $R + r_{\text{cell}}$ of the hint and hands the solver a fragment that contains only those cells’ quads/codes/stars. For an IMX455-class FOV with a tight pointing prior the hint can be the truth pointing $\pm$ a few arcminutes; for a coarse all-sky-mode solve, R is the full π radians and the bundle reader pages through all 3.1 M cells via the LRU-bounded mmap cache.

Because both bundle-side caches—the per-cell shard cache in `Zdc1Bundle` and the per-entry mmap cache in `FsAccessor`—use bounded `LruCaches`, repeated queries against the same bundle do not leak memory: a 1000-case sweep at $R = 4^\circ$ holds approximately $\min(\text{workers} \times \text{cells/region}, \text{cache cap})$ entries resident, which on the test bench peaks around 4 GB regardless of sweep length.

7.6 Empirical Performance vs. Pointing-Hint Radius

We characterise the speed/accuracy trade-off as a function of hint radius using the 1000-case `set2-dr3-mag19` corpus (described in §4.1, regenerated against the depth-9 bundle). Each case is solved with the bench’s regional + scale-hinted configuration ($\pm 25\%$ scale tolerance) at four radii: 0.5° , 1.0° , 2.0° , 4.0° . The cells loaded per query scale as $\pi R^2 \cdot 76.3$ (cells per square degree at $d_b=9$), giving roughly 60, 240, 960, and 3800 cells per region respectively.

Solve rate: 100% across all four radii. The findability fix from §7.4 closes the depth-8 failure mode entirely; every test case in the corpus passes the `log_odds_accept = 20` threshold at every hint radius tested.

Solve time scales near-linearly with R . Median solve time is 1.0 ms at $R = 0.5^\circ$, rising to 12.85 ms at $R = 4^\circ$ (Fig. 5). The median tracks the radius rather than the area because the solver locks onto the correct hypothesis in its first few quad-code matches; what scales with the loaded cell count is the long tail (p99 grows from 18 ms to 563 ms, a $30\times$ spread).

Region-load time scales as R^2 . Median load grows from 22 ms at $R = 0.5^\circ$ to 2.5 s at $R = 4^\circ$ (Fig. 6); the p99 tail balloons from 430 ms to 34 s, dominated by dense Galactic-plane regions whose per-cell shards are $5\text{--}10\times$ larger than the average. Load time is a direct function of cells touched and is the dominant cost at large hints—above $R \approx 1^\circ$ the




Figure 5: Solve-time CDF at four pointing-hint radii against the depth-9 bundle. Each curve is 1000 cases. p50 and p99 callouts shown; the count of approximate cells loaded per query appears in the legend.




Figure 6: Region-load-time CDF (the cost of `bundle.load_region(®ion)` per case). Scales as R^2 because loaded cell count scales as R^2 ; p99 tail is dominated by dense Galactic-plane regions.

bundle read overtakes the solve itself.

Accuracy: median 4 milliarcsec, with a tail at large R . The angular-distance error between the solver’s WCS centre and the truth pointing has a median of $0.004''$ at every hint radius (Fig. 7)—when the solver finds the right hypothesis, accuracy is set by Gaia catalog precision and the WCS-fit residual rather than by the size of the search region. The *tail*, however, splits dramatically: at $R \leq 1^\circ$ the p99 error stays at $0.34''$, but at $R = 2^\circ$ it jumps to $3557''$ and at $R = 4^\circ$ to $9722''$ ($\approx 2.7^\circ$). These are not solver “failures” in the traditional sense—they pass the log-odds verification threshold—they are false-positive WCSes drawn from elsewhere in the loaded region. Counted as wrong-solves above $1''$, the rates are $0.5\% / 0.5\% / 1.5\% / 2.7\%$ across the four radii.

The signed (RA, Dec) residuals for the correct cluster (Fig. 8) are well-centred on zero ($|\text{mean}| < 2$ mas) with $\sigma_{\Delta\alpha \cos \delta} = 32$ mas and $\sigma_{\Delta\delta} = 46$ mas. The slight




Figure 7: Solve-accuracy CDF. Median accuracy is hint-radius-independent at ~ 4 mas; the p99 tail diverges sharply at $R \geq 2^\circ$, populated by false-positive WCSes that nonetheless pass the log-odds threshold. The vertical dashed line is our $1''$ wrong-solve threshold.




Figure 8: Histograms of the signed RA and Dec residual components for the 995 correct cases at $R = 1^\circ$ (error $\leq 1''$). Both centred at zero to within 2 mas; $\sigma_{\Delta\alpha \cos \delta} = 32$ mas, $\sigma_{\Delta\delta} = 46$ mas. Larger Dec scatter reflects TAN-WCS fitting geometry.

asymmetry between the two axes is the characteristic signature of TAN-WCS least-squares fitting: the dec axis has tighter geometric coupling than the $\cos\delta$ -scaled RA component.

Wrong-solve geography. The wrong-solve cases are distributed across the sky but biased toward low-quad-density Galactic-pole regions (Fig. 9), which is consistent with the failure mechanism: at high $|b|$, the bundle holds fewer correct quads in the loaded region, so the relative weight of code-tolerance-passing spurious matches goes up.

7.7 Practical Implications

Operationally, the bundle format reduces the all-sky deep ($G \leq 20$) catalog from a multi-tens-of-GB single index to a 171 GB 2D-sharded structure that is mmap-friendly, queryable in milliseconds-to-seconds depending on the hint, and survives long build pipelines on commodity hardware. The hint-




Figure 9: Wrong-solve case locations (error $> 1''$) for all four hint radii against the depth-9 bundle. Background greyscale shows quads-per-cell density. Wrong-solve concentrations near the Galactic poles trace the catalog’s lowest-density regions.

radius/solve-time trade space is a useful operational dial: pipelines with a ground truth pointing within $\pm 30'$ should expect 1–2 ms median solves and zero false positives, while wider-area surveys with only degree-scale priors should accept the sub-percent false-positive rate or layer a post-solve sanity check (e.g., requiring residual RMS $\leq$ catalog precision) on top of the log-odds verification.

8 Proper Motion at the Observation Epoch

The bundle results in §7.6 were measured against the `set2-dr3-mag19` corpus, which is synthesised at Gaia DR3’s reference epoch ($J2016.0$). For survey pipelines operating close to that epoch the simplification is harmless: stellar drift over a few years stays well below a typical 5 px match radius. Out-of-epoch plates—most acutely *archival* Hubble Space Telescope frames spanning ~ 1990 –2015, but also any pipeline asked to register an image to a time-of-flight epoch—break this assumption. A nominally 500 mas yr $^{-1}$ star drifts $7''$ in 14 years; on a Hubble ACS $0.05''/\text{px}$ plate that is 140 px and the verifier will not associate the detection with the catalog entry at all.

Zodiacal handles this by carrying per-star proper motion through the index and applying it at both read sites that turn (α, δ) into a sky vector: the quad-to-WCS fit (§2.6) and the Bayesian verification pass (§2.7). The on-disk `.zga` per-cell Gaia shard (§7.1) already carries the full DR3 astrometric solution (`pmra`, `pmdec`, reference epoch); enabling this path required only widening the in-memory `IndexStar` record from 4 to 7 fields and threading an `Option<f64>` observa-

tion epoch through the solver configuration.

8.1 Linear Small-Angle Extrapolation

The extrapolation is the standard linear form, applied independently to each axis:

$$\alpha(t_{\text{obs}}) = \alpha(t_{\text{ref}}) + \mu_{\alpha^*} \cdot \frac{t_{\text{obs}} - t_{\text{ref}}}{\cos \delta(t_{\text{ref}})}, \quad (24)$$

$$\delta(t_{\text{obs}}) = \delta(t_{\text{ref}}) + \mu_{\delta} \cdot (t_{\text{obs}} - t_{\text{ref}}), \quad (25)$$

where $\mu_{\alpha^*} \equiv \mu_{\alpha} \cos \delta$ is the Gaia true-sky RA proper motion in mas yr^{-1} , the time difference is in Julian decimal years, and the inverse $\cos \delta$ converts the true-sky rate into a coordinate rate. For fast-moving stars ($\mu \sim 5,000 \text{ mas yr}^{-1}$) over decade-scale gaps this is accurate to $\sim 10 \mu\text{as}$ —two orders of magnitude below any pixel tolerance we care about. Parallax, light-time, and stellar aberration are not applied at the solver level; the `refinement` module’s apparent-place pipeline picks those up for follow-up astrometry once a WCS is in hand.

Two practical guards keep this robust against legacy data. Catalog entries with NaN proper motion—Gaia DR3’s 2-parameter solutions, and all stars in pre-bundle `.zdc1 v1/v2` indexes—are returned unchanged, so a missing measurement degrades silently to the identity rather than corrupting the position with NaN propagation. The observation epoch is itself optional: passing `None` bypasses (24)–(25) entirely and the verifier behaves bit-identically to the pre-PM solver. Synthetic Gaia-epoch corpora therefore continue to produce the same outcomes as §7.6.

8.2 Recovering a High-Proper-Motion Star

A regression test in `src/verify.rs` pins the behaviour with a worked instance. A synthetic detection is placed at pixel (300, 400) under a $1''/\text{px}$ WCS, and the catalog entry is back-extrapolated from a 14-year-distant observation epoch using a 500 mas yr^{-1} proper motion. The forward drift between epochs is $7000 \text{ mas} = 7 \text{ px}$, outside the default 5 px match radius:

- **Without `obs_epoch`:** the verifier projects the index’s 2016 position, lands $\sim 7 \text{ px}$ from the detection, and the star is unmatched.
- **With `obs_epoch`:** the projection first applies Eqs. 24–25 to roll the position forward to the observation epoch, lands inside the match circle, and the star contributes its $\Delta_i \gtrsim 0$ to the log-odds total.

The same wiring is used by `zodiacal-tools bench-bundle`, which gained an `--obs-epoch` flag and an automatic Hubble path that fills `obs_epoch` from the test case’s `t_min_mjd`/`t_max_mjd` midpoint. HST plates therefore solve without operator intervention against a Gaia-epoch bundle, and the bundle on disk gains no new fields.

8.3 When Proper Motion Matters in Practice

The pixel-scale drift admissible without proper-motion correction is

$$\Delta t \lesssim \frac{r_{\text{match}} \cdot s_{\text{px}}}{\mu_{\text{cap}}}, \quad (26)$$

where r_{match} is the verifier match radius (default 5 px), s_{px} is the plate scale, and μ_{cap} is the proper motion above which a star contributes to verification failures. At the bench’s $0.1296''/\text{px}$ and the $\sim 1\%$ of Gaia DR3 stars with $\mu > 100 \text{ mas yr}^{-1}$, the budget is ~ 6.5 years before high-PM stars start dropping out of the match circle. At HST/ACS’s $0.05''/\text{px}$ it shrinks to ~ 2.5 years. Pipelines registering images more than this far from the catalog epoch should set `obs_epoch`; pipelines within the budget see no behavioural change either way.

9 Conclusion

We have presented Zodiacal, a pure-Rust blind astrometry plate solver that achieves 98.3% solve rates with a median time of 1.42 s on simulated narrow-field imagery. Our systematic evaluation of index sharding strategies demonstrates that partitioning the angular scale space into narrow bands is a first-order design decision: moving from 4 coarse bands to 12 fine bands ($\sqrt{2}$ scale factor) improves the solve rate by 2.3 percentage points while maintaining sub-second median solve times for the majority of fields.

The performance benefit arises from two mechanisms: smaller per-band k -d trees that reduce false-positive code matches, and per-index scale filtering that prunes irrelevant bands entirely. These combine to reduce the effective search space by an order of magnitude for typical fields.

The residual 1.7% failure population is characterized by large-backbone geometries that produce code-space collisions, representing a search budget challenge for geometric hashing approaches that narrower bands alone cannot fully resolve.

The legacy single-index design also bounds catalog depth to what fits in working RAM. §7 lifts that limit

by adding a second sharding axis: the `.zdc1.bundle` format pairs per-cell HEALPix shards with the existing scale-band split, and the bundle reader loads only the cells overlapping a user-supplied sky region. Against a depth-9, $G \leq 20$ bundle ($\sim 9 \times 10^8$ stars, 1.6×10^9 quads) and the `set2-dr3-mag19` 1000-case corpus, this delivers a 100% solve rate, 1–13 ms median solve time depending on hint radius, and median 4 milliarcsec positional accuracy. We characterise the speed/accuracy trade-off across hint radii of 0.5° – 4.0° , including the false-positive WCS regime that emerges at large radii where multiple quads can plausibly verify the same field.

Out-of-epoch plates are handled by §8: per-star proper motion is carried through the bundle’s `.zga` shards and applied at WCS fit and verification when the caller supplies an observation epoch. The solver continues to behave bit-identically to the pre-PM path when no epoch is given, so synthetic Gaia-epoch corpora are unaffected; archival HST plates now solve against the same depth-9 bundle without operator intervention.

Zodiacal is open-source and available at <https://github.com/OrbitalCommons/zodiacal>.

Acknowledgments

This work makes use of data from the European Space Agency mission Gaia (<https://www.cosmos.esa.int/gaia>), processed by the Gaia Data Processing and Analysis Consortium.

References

- Matthew Goodman. starfield: Astronomical data reduction toolkit, 2025. URL <https://github.com/OrbitalCommons/starfield>. Rust crate for star catalogs, coordinate systems, and source extraction.
- K. M. Górski, E. Hivon, A. J. Banday, B. D. Wandelt, F. K. Hansen, M. Reinecke, and M. Bartelmann. HEALPix: A Framework for High-Resolution Discretization and Fast Analysis of Data Distributed on the Sphere. *The Astrophysical Journal*, 622(2): 759–771, 2005. doi: 10.1086/427976.
- Dustin Lang, David W. Hogg, Keir Mierle, Michael Blanton, and Sam Roweis. Astrometry.net: Blind Astrometric Calibration of Arbitrary Astronomical Images. *The Astronomical Journal*, 139(5):1782–1800, 2010. doi: 10.1088/0004-6256/139/5/1782.
- D. T. Lee and C. K. Wong. Worst-case analysis for region and partial region searches in multidimensional binary search trees and balanced quad trees. *Acta Informatica*, 9(1):23–29, 1977. doi: 10.1007/BF00263763.